

Graph Representation Learning, Relational GCNs, and Knowledge Graph Inference

Adam Abramowitz Max Desantis Isaac Dreeben Abhi Mummaneni

Abstract

Much of modern data is relational — social networks, knowledge graphs, molecular structures, and narrative scenes all encode meaning through connections, not just attributes. This paper surveys the landscape of graph representation learning, tracing the development from classical spectral methods through graph convolutional networks, relational GCNs, and generative graph models. We then ground these ideas in two experiments using the Relational Graph Convolutional Network (R-GCN). First, we validate the architecture on a Wikidata temporal knowledge graph, demonstrating that relational message passing improves link prediction over a shallow baseline across all metrics. Second, we construct a novel dataset of over 500 murder mystery plots, using an LLM to extract relational knowledge graphs, and train the R-GCN to identify villains from circumstantial evidence alone. The model achieves 90.4% accuracy and 0.762 villain F_1 , with a +14-point precision advantage over a feature-only logistic regression baseline — a gap attributable specifically to typed relational structure rather than graph topology. The model further generalizes to a real unsolved case, producing suspect rankings consistent with current forensic evidence. Together, these results demonstrate that capturing typed relational structure leads to more effective learning across a diverse range of graph-structured problems.

Contents

1	Introduction	3
2	Structured Graph Representation Learning	3
2.1	From Spectral Methods to Graph Convolutional Networks	3
2.1.1	Early Methods in Graph Learning	3
2.1.2	Graph Convolutional Networks	4
2.1.3	Graph Autoencoders and Generative Graph Models	4
2.1.4	Relational Graphs and Relational GCNs	5
2.1.5	Message Passing, Attention, and Graph Transformers	6
2.1.6	Expressivity and Limitations	7
2.2	Deeper Autoencoders: VAEs and Adversarial Autoencoders	7
2.3	Generative Adversarial Networks	7
2.3.1	Adversarial Autoencoders	8
2.3.2	Adversarially Regularized Graph Autoencoders	8
3	Experiments: R-GCN Deep Dive	9
3.1	Wikidata Knowledge Graph	9
3.1.1	Results	10
3.1.2	Discussion	10
3.2	Detective Knowledge Graph	11
3.2.1	Task and Motivation	11
3.2.2	Dataset Construction	12
3.2.3	Model and Training	12
3.2.4	Restricting the Objective to Villain Identification	13
3.2.5	Results	13
3.2.6	Discussion	15
4	Conclusion	16

1 Introduction

Many real-world problems involve modeling highly structured graph data: social networks, knowledge graphs, time-series data, molecular structure, and much more. The challenge is to build models that capture deep structures beyond just surface level connections. We explore this challenge through both theory and practice. This paper will be divided into two parts.

First we'll provide an expository survey of the modern techniques and models for embedding highly structured graphs for downstream tasks like deep learning and reconstruction. We'll trace the development of graph convolutional networks starting with the spectral methods we saw in class. We'll then see how these methods can be transformed into deeper generative models with richer latent spaces.

The second part of the paper will be a deep dive into modeling relational graphs via the Relation Graph Convolutional Network introduced by Schlichtkrull et al. ADD CITATION HERE. We provide two experiments, both embedding relational graphs with the R-GCN architecture. The first graph is a Wikidata knowledge graph which we embed to demonstrate the effectiveness of embedding relational structure in addition to purely topological structure. Next, we use an LLM to generate knowledge graphs of over 500 detective story plots. We then embed these graphs with the R-GCN with the goal of predicting villains of the murder mysteries. The model then generalizes to predict the real-world Zodiac killer suspect. These experiments will show that capturing deeper structure leads to more effective learning across a diverse class of data types and settings.

2 Structured Graph Representation Learning

2.1 From Spectral Methods to Graph Convolutional Networks

2.1.1 Early Methods in Graph Learning

Early graph-learning methods were motivated by the observation that graph structure can be used as an inductive bias for representation learning. In Laplacian Eigenmaps, Belkin and Niyogi Belkin and Niyogi [2001] construct a weighted graph over data points and seek low-dimensional embeddings that keep neighboring nodes close. Let W be the weighted adjacency matrix, D the diagonal degree matrix, and $L = D - W$ the combinatorial graph Laplacian. The embedding matrix $Y \in \mathbb{R}^{|V| \times d}$ is obtained by solving

$$\min_Y \frac{1}{2} \sum_{i,j} W_{ij} \|Y_i - Y_j\|_2^2 = \min_Y \text{tr}(Y^\top L Y), \quad \text{s.t. } Y^\top D Y = I, Y^\top D \mathbf{1} = 0.$$

This objective makes the geometric assumption explicit: if two examples are adjacent in the graph, their learned representations should vary smoothly across that edge. The solution is given by the generalized eigenvectors of $Ly = \lambda Dy$, connecting graph embeddings to spectral approximations of the Laplace–Beltrami operator on an underlying manifold.

This smoothness viewpoint was further developed in manifold regularization Belkin et al. [2006], where the graph is used to regularize functions defined on data points. For a graph signal $f : V \rightarrow \mathbb{R}$, the graph smoothness penalty is

$$\mathcal{R}_G(f) = \frac{1}{2} \sum_{i,j} W_{ij} (f_i - f_j)^2 = f^\top L f.$$

Thus, learning on graphs can be interpreted as learning functions whose values change slowly across high-weight edges. Shuman et al. Shuman et al. [2013] generalized this interpretation through the

language of graph signal processing. If $L = U\Lambda U^\top$ is an eigendecomposition of the graph Laplacian, then the graph Fourier transform of a signal $x \in \mathbb{R}^{|V|}$ is

$$\hat{x} = U^\top x, \quad x = U\hat{x}.$$

Eigenvectors associated with small eigenvalues correspond to smooth, low-frequency variation over the graph, while larger eigenvalues correspond to high-frequency oscillations. This spectral viewpoint provides the foundation for defining graph convolutional filters.

2.1.2 Graph Convolutional Networks

Spectral graph neural networks extend convolution from Euclidean domains to irregular graph domains by filtering graph signals in the Laplacian eigenbasis Bruna et al. [2014], Defferrard et al. [2016], Kipf and Welling [2017]. Given a signal x and a spectral filter g_θ , convolution can be written as

$$g_\theta * x = U g_\theta(\Lambda) U^\top x,$$

where $g_\theta(\Lambda)$ scales each graph-frequency component independently. This formulation is expressive but depends on the eigendecomposition of the graph Laplacian, which is expensive for large graphs and not naturally transferable across graphs with different eigenbases.

Chebyshev spectral networks address this limitation by approximating the filter as a polynomial of the rescaled Laplacian Defferrard et al. [2016]:

$$g_\theta * x \approx \sum_{k=0}^K \theta_k T_k(\tilde{L})x, \quad \tilde{L} = \frac{2}{\lambda_{\max}} L - I,$$

where T_k denotes the k -th Chebyshev polynomial. Because this filter is a K -th order polynomial in L , it is localized to K -hop neighborhoods and avoids explicit eigendecomposition.

Kipf and Welling [2017] further simplified this construction by using a first-order approximation and a renormalized adjacency matrix. With self-loops $\tilde{A} = A + I$ and degree matrix $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$, a standard GCN layer is

$$H^{(\ell+1)} = \sigma\left(\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} H^{(\ell)} W^{(\ell)}\right).$$

This layer averages transformed node features over local neighborhoods, making GCNs an efficient and widely used architecture for semi-supervised node classification. However, repeated applications of the propagation operator can also cause node representations to become indistinguishable. Li, Han, and Wu Li et al. [2018] interpret GCN propagation as Laplacian smoothing; after K propagation steps, the representation is dominated by powers of the normalized adjacency:

$$H^{(K)} \approx \left(\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}\right)^K X.$$

This smoothing perspective motivates later work on deeper GNNs, residual connections, attention mechanisms, and architectures designed to preserve long-range or high-frequency information.

2.1.3 Graph Autoencoders and Generative Graph Models

A second line of work relevant to graph representation learning comes from autoencoders and variational autoencoders. The variational autoencoder framework of Kingma and Welling Kingma

and Welling [2014] introduces a latent-variable model in which an encoder approximates the posterior distribution over latent variables and a decoder reconstructs the observed data. For data x , latent variable z , prior $p(z)$, encoder $q_\phi(z | x)$, and decoder $p_\theta(x | z)$, the evidence lower bound is

$$\mathcal{L}_{\text{VAE}}(\theta, \phi; x) = \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x | z)] - \text{KL}(q_\phi(z | x) \parallel p(z)).$$

The first term encourages accurate reconstruction, while the KL term regularizes the latent space toward the prior. This framework naturally supports probabilistic generation by sampling from $p(z)$ and decoding.

Variational Graph Autoencoders (VGAEs) adapt this idea to graph-structured data by using a GCN encoder and an edge-probability decoder Kipf and Welling [2016]. Given node features X and adjacency matrix A , the encoder factorizes over nodes as

$$q_\phi(Z | X, A) = \prod_{i=1}^{|V|} q_\phi(z_i | X, A), \quad q_\phi(z_i | X, A) = \mathcal{N}(z_i | \mu_i, \text{diag}(\sigma_i^2)),$$

and the decoder predicts edges through an inner product:

$$p_\theta(A_{ij} = 1 | z_i, z_j) = \sigma(z_i^\top z_j).$$

This model connects graph representation learning with link prediction: the latent embedding must encode enough structural information to reconstruct the observed adjacency matrix.

Subsequent generative graph models extend VGAEs from link prediction toward whole-graph generation. GraphVAE Simonovsky and Komodakis [2018] decodes latent variables into probabilistic node and edge tensors for small graphs, while Graphite Grover et al. [2019] improves the decoder by iteratively refining a generated graph. These models can be viewed as learning a distribution over graphs,

$$p_\theta(G) = \int p_\theta(G | z) p(z) dz,$$

where the decoder must capture both local edge dependencies and global graph-level constraints. This generative perspective complements discriminative GNNs by emphasizing reconstruction, latent structure, and graph synthesis.

2.1.4 Relational Graphs and Relational GCNs

Many real-world graphs are not homogeneous: edges may represent different relation types, and nodes may belong to different semantic categories. Relational GCNs extend the GCN update rule to multi-relational graphs by assigning relation-specific transformations Schlichtkrull et al. [2018]. For relation set $\mathcal{R}$, relation-specific neighborhood $\mathcal{N}_i^r$, and normalization constant $c_{i,r}$, the R-GCN update is

$$h_i^{(\ell+1)} = \sigma \left(\sum_{r \in \mathcal{R}} \sum_{j \in \mathcal{N}_i^r} \frac{1}{c_{i,r}} W_r^{(\ell)} h_j^{(\ell)} + W_0^{(\ell)} h_i^{(\ell)} \right).$$

This formulation separates messages by relation type, making it suitable for knowledge graphs and other structured domains in which edge labels carry semantic meaning.

Heterogeneous graph models further generalize this idea by allowing both node-type-specific and edge-type-specific transformations. The Heterogeneous Graph Transformer of Hu et al. Hu et al.

[2020] uses attention mechanisms whose parameters depend on the source node type, target node type, and relation type. A simplified relation-aware attention score can be written as

$$\alpha_{ij}^r = \text{softmax}_{j \in \mathcal{N}_i^r} \left(\frac{(Q_{\tau(i)} h_i)^\top R_r (K_{\tau(j)} h_j)}{\sqrt{d}} \right),$$

where $\tau(i)$ denotes the type of node i and r denotes the relation connecting nodes i and j . This line of work shows how similarity graphs used in early spectral methods can be extended to typed, relational, and semantically structured graphs.

2.1.5 Message Passing, Attention, and Graph Transformers

Message passing provides a unifying abstraction for many GNN architectures. Gilmer et al. [2017] formulate neural message passing in terms of a message function, aggregation operation, and update function. At layer t , a generic message passing update is

$$m_i^{(t+1)} = \sum_{j \in \mathcal{N}(i)} M_t(h_i^{(t)}, h_j^{(t)}, e_{ij}), \quad h_i^{(t+1)} = U_t(h_i^{(t)}, m_i^{(t+1)}).$$

GCNs, R-GCNs, and many spectral approximations can be interpreted as instances of this template with different choices of message function, normalization, and aggregation.

Graph Attention Networks make the connection between message passing and attention explicit by replacing fixed neighborhood averaging with learned attention weights Veličković et al. [2018]. For neighboring nodes i and j , attention coefficients are computed as

$$\alpha_{ij} = \text{softmax}_{j \in \mathcal{N}(i)} \left(\text{LeakyReLU} \left(a^\top [W h_i \parallel W h_j] \right) \right),$$

and the node update becomes

$$h'_i = \sigma \left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij} W h_j \right).$$

Attention allows the model to assign different weights to different neighbors, rather than assuming that all adjacent nodes contribute equally.

Transformer-style graph models generalize this principle by combining global or constrained attention with structural encodings. Dwivedi and Bresson [2020] adapt Transformers to arbitrary graphs using graph-aware positional encodings, while Graphormer Ying et al. [2021] demonstrates that full attention can be effective when enriched with structural biases such as shortest-path distances and edge encodings. A common way to express graph-biased attention is

$$\text{Attn}(Q, K, V; B) = \text{softmax} \left(\frac{QK^\top}{\sqrt{d}} + B \right) V,$$

where B_{ij} encodes graph structure. In constrained attention, one may set

$$B_{ij} = \begin{cases} 0, & (i, j) \in E \text{ or } i = j, \\ -\infty, & \text{otherwise,} \end{cases}$$

which recovers neighborhood-restricted attention; richer choices of B permit long-range interactions while preserving graph inductive bias. Recent graph language models Plenz and Frank [2024] further connect this line of work to language modeling by incorporating graph structure into pretrained, Transformer-like architectures.

2.1.6 Expressivity and Limitations

Although message passing architectures are broadly useful, their expressivity is limited. Xu et al. [2019] relate the discriminative power of GNNs to the Weisfeiler–Lehman (WL) graph isomorphism test. In WL-style refinement, node colors are updated by hashing the previous color together with the multiset of neighboring colors:

$$c_i^{(k)} = \text{HASH}\left(c_i^{(k-1)}, \left\{c_j^{(k-1)} : j \in \mathcal{N}(i)\right\}\right).$$

Similarly, standard message passing GNNs update node embeddings by aggregating multisets of neighbor representations. As a result, many GNNs cannot distinguish graph structures that the 1-WL test cannot distinguish.

A separate limitation is over-squashing: information from a large receptive field must be compressed into fixed-dimensional node embeddings. After T message passing layers, a node representation can depend on all nodes within T hops,

$$h_i^{(T)} = \Phi_T(\{x_j : \text{dist}(i, j) \leq T\}) \in \mathbb{R}^d,$$

but the dimension d remains fixed even when the number of contributing nodes grows exponentially with distance. Alon and Yahav [2021] identify this bottleneck as a practical obstacle for long-range reasoning on graphs.

These limitations motivate architectures that augment message passing with virtual edges, positional encodings, higher-order structures, or global attention. However, as seen in Veličković [2022] many approaches that appear to go beyond message passing can be reframed as message passing over an augmented graph,

$$G' = (V', E'), \quad E \subseteq E',$$

where added edges, nodes, or features change the computational graph over which information flows. In conclusion, though the many GCN architectures and graph learning methods seem distinct, each can be understood in the context of the Laplacian of the graph and distinguished by their choices of graph structure, message aggregation, and inductive bias.

2.2 Deeper Autoencoders: VAEs and Adversarial Autoencoders

2.3 Generative Adversarial Networks

Goodfellow et al. introduced Generative Adversarial Nets, which optimize two functions G and D through the following min-max objective:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))].$$

D must discriminate between the prior data distribution p_{data} , and the generator’s output $G(z)$, with $z \sim p_z$.

The resultant Nash equilibrium is that G learns to model p_{data} , and D cannot distinguish between generated samples and samples from p_{data} . Note that this only requires the ability to sample from the prior distribution. This enables a generalization of VAEs towards more complex and hard to define latent regularizations. It is also important to note that, as Goodfellow et al. mention, memorization is less likely in a GAN as information from the prior distribution does not flow directly into the generator.

2.3.1 Adversarial Autoencoders

Makhzani et al. connect GANs and autoencoders by using GANs to regularize the latent space. A traditional variational autoencoder looks to minimize the following:

$$\mathcal{L}_{\text{VAE}}(\theta, \phi; x) = -\mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x | z)] + \text{KL}(q_\phi(z | x) \parallel p(z)).$$

This however requires being able to form

$$q_\theta(z | x)$$

which is not a probability distribution in a deterministic encoder. Consider a standard autoencoder. How would you produce a distribution over latents for a given data point? Adversarial Autoencoders introduce the aggregated latent, which is a distribution over all of the latent variables, not just the latent distribution for a given data point.

$$q(z) = \int_x q(z | x) p_d(x) dx$$

Then the autoencoder is trained using alternating minimization, where first the autoencoder is trained to accurately reconstruct the inputs:

$$L = \frac{1}{N} \sum_{i=1}^n (y_i - D(z_i))^2 \quad \text{with} \quad z_i = E(x_i)$$

where E is the encoder, D is the decoder, (x_i, y_i) is the labeled training data. Afterwards, the model is trained to minimize:

$$\min_E \max_H V(H, E) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log H(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - H(E(z)))].$$

where E is the encoder and H is the discriminator.

2.3.2 Adversarially Regularized Graph Autoencoders

Pan et al. combined concepts from VGAEs and AAEs to use adversarial methods to regularize the latent space of graph autoencoders. They use a Graph Convolutional Network for their encoding model, and

$$p(\hat{A}_{ij} = 1 | z_i, z_j) = \sigma(z_i^\top z_j)$$

as their decoder for link prediction.

The paper used two concepts from adversarial autoencoders: the aggregated latent and the alternating minimization scheme between minimization reconstruction loss and adversarially regularize the latent space. For reconstruction loss minimization, the authors propose using the log-likelihood

$$L_0 = \mathbb{E}_{q(Z|X,A)} [\log p(\hat{A} | Z)]$$

for deterministic models, but also propose using

$$L_1 = \mathbb{E}_{q(Z|X,A)} [\log p(\hat{A} | Z)] - \text{KL} [q(Z | X, A) \parallel p(Z)]$$

for variational models, thus using both VAE regularization and adversarial regularization.

In both cases, Z are the latent vectors, A is the adjacency matrix, and X are node features.

3 Experiments: R-GCN Deep Dive

We now turn from our exposition on structured graph learning to two applications of the relational graph convolutional network (R-GCN) model introduced by Schlichtkrull et al. [?] ADD CITATION HERE. First, we embed a wikidata temporal knowledge graph with the goal of link prediction. We chose this dataset to demonstrate that the R-GCN message passing technique improves predictive capabilities over naive models that don't see the relational structure. Secondly, we test the R-GCN encoder on an entity classification task with a knowledge graph built from detective stories in movies, novels, TV shows, etc.

3.1 Wikidata Knowledge Graph

Before applying the R-GCN to the detective corpus, we first validate it on tkgl-smallpedia, a temporal knowledge graph benchmark drawn from Wikidata [?] CITATION HERE. The dataset consists of ordered quadruples (s, r, o, t) encoding that entity s is related to entity o by relation r at time t . Entities are Wikidata items (persons, cities, institutions, awards, etc.) and relations are Wikidata properties (member of sports team, country of citizenship, award received, doctoral student, etc.). Table 1 provides some dataset statistics and Figure 1 shows a subgraph of some of the nodes connected to the Kingdom of Italy for a visual example.




Figure 1: Subgraph of the tkgl-smallpedia dataset. The red edges have the relation *Country of Citizenship*, the blue edges have the relation *Country*, and the purple edges have the relation *located in the administrative territorial entity*.

For the purpose of this experiment, we will ignore the time dependence term t for each edge and view our task as static link prediction: given any two entries of the triple, (s, r, o) , predict the third, among all entities (or relations).

The purpose of this experiment is specifically to isolate the contribution of relational message passing. To that end, we compare two models that share an identical DistMult decoder [?] but differ in how entity representations are produced:

- **Shallow DistMult.** Each entity receives a learned embedding with no graph encoder. The decoder scores a triple with the scoring function $f(s, o, r) = h_s^\top R_r h_o$. This baseline captures

co-occurrence statistics but ignores the graph structure entirely.

- **R-GCN + DistMult.** First, we embed our feature matrix X^0 with a two-layer R-GCN encoder (hidden dimension 64, 30 basis matrices, dropout 0.1). Then train on the same DistMult scoring and loss function.

Both models are trained for 50 epochs with the Adam optimiser (learning rate 10^{-3} , batch size 4,096) using binary cross-entropy loss against one randomly corrupted negative per positive triple. Evaluation follows the filtered ranking protocol with 50 negatives per positive on the held-out test split.

Property	Value	Split	Edges
Temporal edges	1,100,752	Training	775,514
Unique nodes	47,433	Validation	162,066
Relation types	566	Test	163,172
Time span	1900–2024		

Table 1: Dataset statistics for tkgl-smallpedia.

3.1.1 Results

Table 2 reports Mean Reciprocal Rank (MRR) and Hits@ K for $K \in \{1, 3, 10\}$ on the test set.

Model	MRR	Hits@1	Hits@3	Hits@10
Shallow DistMult	0.2238	0.1739	0.2084	0.2676
R-GCN + DistMult	0.2900	0.2026	0.2690	0.4698

Table 2: Link prediction on TKGL-SMALLPEDIA (test set, 50 negatives per positive, filtered ranking). Best result per metric in **bold**.

The R-GCN encoder improves over the shallow baseline across all four metrics, with relative gains of +29.6% in MRR, +16.5% in Hits@1, +29.1% in Hits@3, and +75.6% in Hits@10.

3.1.2 Discussion

State-of-the-art link prediction methods on knowledge graphs have Hits@ K hovering around 0.6. So, our R-GCN encoder is not quite there. However, we do see that the R-GCN encoder does strictly improve every metric across the board compared to the unstructured shallow DistMult baseline. The most striking result is the Hits@10 improvement: the R-GCN places the correct object in the top 10 candidates for nearly half of all test queries (0.4698), compared to only one quarter for the shallow baseline (0.2676). Some Hits@1 and Hits@3 but large gains with Hits@10 is consistent with what one would expect from a structural encoder. The R-GCN does not always rank the single correct answer first, but it reliably pulls it into the set of plausible candidates by recognizing that entities sharing similar relational neighborhoods are likely to participate in similar triples. The shallow model, lacking any notion of neighborhood, relies entirely on direct co-occurrence statistics and so it struggles to generalize to less-frequent entities or relation types.

The small, but still positive, MRR and Hits@1 gains (+29.6% and +16.5% respectively) may reflect the difficulty of precise top-1 ranking on a knowledge graph with ambiguous entities. The

asymmetry between Hits@1 and Hits@10 suggests there is still room for improvement in the encoder. For instance, we could implement tighter hyperparameter tuning, attention-weighted message passing (cf. Graph Attention Networks [Veličković et al., 2018]), CITATION HERE or we could have incorporated the temporal index t that we set aside for this experiment. One additional shortfall of the DistMult scoring function $f(s, r, o) = h_s^\top R_r h_o$ is that it is bidirectional. The matrix R_r is diagonal which means $h_s^\top$ and h_o commute through R_r , i.e., $h_s^\top R_r h_o = h_o^\top R_r h_s$. So, the model scores “ A is related to B ” the same as “ B is related to A .” Adding synthetic triples (o, r^{-1}, s) with reverse edges would solve this issue. We implement this change to address the obvious bidirectional issue with the murder mystery experiment. But, implementing this change for tkg1-smallpedia may also improve the DistMult scoring metrics.

Just to see how the scoring function operates, if we ask to predict the object in the incomplete triple (Marie Curie, Winner, ?), the top 5 predicted objects with their scores are shown in Table 3.

Rank	Predicted Entity	Score
1	Nobel Prize in Chemistry [†]	5.0768
2	Eddington Medal	4.3701
3	Nobel Prize in Physics [†]	3.5951
4	National Medal of Science	3.5567
5	Henry Draper Medal	3.4259

Table 3: Top-5 predicted completions for the query (Marie Curie, winner, ?). Factually correct answers are marked with [†].

Clearly every prediction is a real, prestigious scientific award. So, we see that the model has learned that Marie Curie should be clustered with the highly accomplished scientists. Additionally, the model did train on the triple (Marie Curie, award received, Nobel Prize in Chemistry). However, the “Marie Curie” node did not have the word “winner” as one of its training edges. This suggests the model has generalized across semantically related but structurally distinct relation types.

Taken together, these results validate the hypothesis that explicitly encoding relation type during neighborhood aggregation produces strictly better representations than a flat embedding table, even on a clean, well-curated benchmark where co-occurrence alone provides a strong signal. In Section 3.2 we test this same architecture on a noisier graph extracted from narrative text.

3.2 Detective Knowledge Graph

3.2.1 Task and Motivation

To further test relational message passing on a problem that is structurally relational, we built a knowledge-graph dataset of murder mysteries and trained the R-GCN ADD CITATION HERE to play detective. The setup is the following: given a story whose victim is known, the model is shown the full cast of characters, their per-node features (motive, alibi, concealment, etc.), and the typed relationships among characters, locations, occupations, and organizations. It must predict which character is the villain. Crime-revealing edges (**kills**, **kills_inv**, **assaults**, **financial crime**) are masked at evaluation time, so the model has to reason from circumstantial evidence: motive, alibi, concealment, spatial co-presence, and social ties. The main result is that the R-GCN reaches 90.4% overall accuracy and 0.762 villain F_1 on a 5-seed cross-validation, with a +14-point villain-precision advantage over a logistic-regression baseline that uses the same node features without graph structure.

3.2.2 Dataset Construction

We assembled a master list of 755 candidate murder mysteries spanning novels, short stories, films, TV episodes, and podcasts. After web scraping, deduplication, and quality filtering, the working set is 576 stories (340 novels, 126 TV episodes, 62 films, 28 podcasts, 20 short stories). A custom scraper pulled each story’s synopsis. Synopses shorter than 150 words or scoring below 0.4 on a four-signal quality filter (word count, named-character count, event-verb density, presence of a resolution sentence) were excluded.

Each synopsis was converted into a relational graph by a two-pass LLM extraction pipeline running locally on `mixtral:8x7b` via Ollama. Pass 1 extracts entities (characters, occupations, locations, organizations) with their feature vectors; pass 2 takes those entities plus the synopsis and extracts edges. The two-pass design substantially improves edge density relative to single-pass extraction, which routinely missed 40–60% of character–character relations. The number of passes was a consequence of the extraction LLM model used and not by design. A sample set was extracted via Anthropic’s Claude api, but proved expensive because of the large number of tokens processed for extracting each relational graph. We subsequently explored several models that could be run locally and settled on `mixtral:8x7b` which produced good results but took about 5 minutes per extraction (roughly 50 hours). A subsequent normalization step collapses roughly 1,500 free-text relation strings into 8 coarse classes, and the loader duplicates each edge (A, r, B) as an inverse (B, r_{inv}, A) , giving 16 relation types in training: `kills`, `harms`, `investigates`, `deceives`, `personal_bond`, `professional`, `spatial`, `social`, and their inverses.

The resulting graph set contains 14,020 nodes and 31,894 edges (25,732 training, 3,166 validation, 2,996 test). The median graph has 9 characters, 24 edges, 5 locations, and 6 occupations. Of 5,907 labeled characters, 49.6% are Uninvolved, 22.5% Villain, 17.7% Victim, 9.1% Witness, and 1.1% unlabeled.

Initial training revealed that roughly 20% of test stories had villains with no extracted feature signal — -1 (UNK) across motive, concealment, and hidden-relationship features. We subsequently pulled new synopses for these stories specifically using number of words as the key metric for sufficiency. Three rounds of targeted re-extraction lifted villain F_1 from 0.703 to 0.762 and reduced the number of test cases neither model could solve from 11 to 0. Four stories were excluded entirely as out-of-scope: two real-world unsolved cases (FLM_047 *Zodiac*, POD_035 *In the Dark: Season 3*, both retained as held-out inference probes since both cases are unresolved) and two with extraction-incompatible villains (TVE_089 *Vera: Hidden Depths*, where the labeled villain was a generic phrase rather than a named character; TVE_093 *Spiral: Series 4*, an organized-crime serial with no single perpetrator). The dataset’s stated scope is therefore single-perpetrator detective fiction.

3.2.3 Model and Training

We trained an R-GCN following Schlichtkrull et al. ADD CITATION HERE. The encoder is a two-layer R-GCN with per-relation neighbor normalization $c_{i,r} = |\mathcal{N}_i^r|$ (rather than the global degree), separate weight matrices W_r for each of the 16 relation types, and a self-loop weight W_0 :

$$h_i^{(\ell+1)} = \sigma \left(\sum_{r \in \mathcal{R}} \sum_{j \in \mathcal{N}_i^r} \frac{1}{c_{i,r}} W_r^{(\ell)} h_j^{(\ell)} + W_0^{(\ell)} h_i^{(\ell)} \right).$$

Hidden dimension is 32 with dropout 0.4. Because each base relation and its inverse are assigned distinct weight matrices, the encoder can learn directional patterns — “ A killed B ” and “ B was

killed by A ” produce different updates. Two decoders sit on top of the shared encoder: a DistMult head for link prediction and a linear classifier for character labels. The joint training objective is

$$\mathcal{L} = \lambda_{\text{link}} \cdot \mathcal{L}_{\text{link}} + \lambda_{\text{cls}} \cdot \mathcal{L}_{\text{cls}},$$

where $\mathcal{L}_{\text{link}}$ is binary cross-entropy over positive edges and corrupted negatives (with subject *or* object randomly replaced) and $\mathcal{L}_{\text{cls}}$ is class-weighted cross-entropy over character labels using inverse-frequency weights. Splits are story-level (462 training / 57 validation / 57 test), so all characters in a held-out story are unseen at evaluation. Training runs for 50 epochs and converges by epoch 5–10; total training time was under ten minutes.

3.2.4 Restricting the Objective to Villain Identification

Initial runs used a four-class classification head (Villain / Victim / Witness / Uninvolved). On a single seed this hit roughly 66% overall accuracy, but cross-class behavior was poor: Witness recall was 7% and Victim recall 30%, because Witness and Victim are minority classes (9% and 18% of labels) with weak features distinguishing them from Uninvolved characters. The detective scenario, however, does not require this distinction — the victim is known by hypothesis, and from the detective’s point of view a Witness and an Uninvolved character are equally “not the perpetrator.”

We therefore reduced the classification to binary Villain vs. Non-Villain identification by mapping $\{\text{Victim, Witness, Uninvolved}\} \rightarrow \text{Non-Villain}$ in $\mathcal{L}_{\text{cls}}$. The encoder still ingests the full graph, all node features, and the (masked) edge structure; only the supervised target changes. This produced two key results. First, overall accuracy rose from $\sim 66\%$ to 90.4%, because the model is no longer penalized for the Witness/Victim/Uninvolved separation it cannot reliably make from the available evidence. Second, the binary loss focuses gradient signal on the one distinction the task actually demands. Villain F_1 rose from 0.71 (four-class) to 0.762 (binary, cross-validated), and the story-level clean-solve rate — correctly identifying at least one villain with no false accusations of innocents — rose from 56% to 77.2%. All precision and recall figures in Section 3.2.5 refer specifically to this binary task with crime edges masked.

3.2.5 Results

Across 5 random story-level splits, the R-GCN reaches 0.825 villain precision, 0.708 villain recall, 0.762 villain F_1 , and 90.4% overall accuracy. A logistic-regression baseline using the same 8 character features without graph structure reaches 0.684 precision, 0.752 recall, 0.715 F_1 , and 86.9% accuracy. Table 4 reports the full results by approach. The R-GCN’s overall accuracy variance across seeds is extremely tight (± 0.007), and its precision advantage over the LogReg (+14.1 points) is consistent across every seed.

Precision/recall split. The LogReg has higher recall (0.752 vs. 0.708) but produces roughly twice the false accusations: 39 false positives versus the R-GCN’s 19 on the seed-42 test set, with proportionally fewer victims and witnesses falsely accused (R-GCN: 10 Uninvolved, 7 Witnesses, 2 Victims; LogReg: 22, 13, 4). The disagreement pattern is asymmetric and consistent across all five seeds: when the R-GCN wins a disagreement it is always by clearing an innocent the LogReg flagged, and when the LogReg wins it is always by catching a villain the R-GCN missed. The graph structure helps the model rule suspects *out* more than it helps find them — precisely the behavior expected from a model that can reason about alibis, social context, and spatial separation. As Table 5 confirms, this asymmetric advantage cannot be reproduced by graph topology alone: a Laplacian-eigenmaps embedding without features collapses to 25.5% accuracy, and combining

Table 4: Villain identification, binary task, crime edges masked, narrative-prominence features excluded. Mean $\pm$ standard deviation over 5 random story-level train/val/test splits (seeds 42, 123, 456, 789, 2024); 462 training stories, 57 test stories per split. Story-level outcomes are reported on the seed-42 split for reference.

Metric	R-GCN	LogReg baseline
Villain Precision	0.825 $\pm$ 0.033	0.684 $\pm$ 0.060
Villain Recall	0.708 $\pm$ 0.019	0.752 $\pm$ 0.046
Villain F_1	0.762 $\pm$ 0.025	0.715 $\pm$ 0.045
Overall Accuracy	0.904 $\pm$ 0.007	0.869 $\pm$ 0.019
At least one villain caught (seed 42)	56/57 (98.2%)	57/57 (100.0%)
Clean solves, no false accusations (seed 42)	44/57 (77.2%)	38/57 (66.7%)
False accusations on test set (seed 42)	19	39

Table 5: Comparison against simpler graph baselines on the same story-level test split. Spectral baselines use a simplified character-only weighted graph with edge weights derived from direct character-character relations and shared locations, organizations, and occupations.

Approach	Test Acc.	V. Prec.	V. Recall	V. F_1
Spectral only (4-dim Laplacian, no features)	25.5%	0.145	0.093	0.113
Features-only LogReg (no graph)	58.5%	0.737	0.676	0.705
Features + spectral (concatenated)	58.8%	0.727	0.667	0.696
R-GCN (typed relational MP)	90.4%	0.825	0.708	0.762

spectral embeddings onto the LogReg adds +0.3 points. The signal is not in the connectivity but in the typed relations.

Held-out case study: the Zodiac murders This late 1960s/early 1970s real-world case does not have a confirmed villain(s) - it is an unsolved case. The 2007 Fincher film, *Zodiac* (FLM 047), covers this case in which Arthur Leigh Allen is the police-named prime suspect. For comparison, we extracted the synopsis (and subsequent graph) using only evidentiary information (police reports, investigative reporting, etc.). ZODIAC_FACTUAL was rebuilt from FBI, SFPD, Vallejo PD, Napa County SO, Solano County SO, and CA DOJ sources, with 2023–2025 forensic updates including the December 2023 DNA identification of victim Donna Lass and the explicit exclusion of Allen as suspect by DNA. We then asked the model to rank suspects by predicted villain probability. The crime edges are masked, and the model has never seen either of these graphs.

Table 6: Suspect rankings on two held-out, never-trained renderings of the Zodiac Killer case. The two synopses describe the same real-world unsolved case with different evidentiary emphases.

Synopsis	#1 suspect	P(v.)	Allen rank	Allen P(v.)	Real-world status of #1
FLM_047 (2007 fiction)	Arthur Leigh Allen	1.0000	#1 of 20	1.0000	Police prime suspect at time of film
ZODIAC_FACTUAL (factual, 2023–2025)	Lawrence Kane	0.9587	#38 of 38	0.0004	Vallejo PD’s officially-developed suspect since 1991; re-centered by Dec. 2023 Donna Lass DNA identification

The two versions yield different prime suspects, and both are defensible by the evidence the model was given. In the fictional film version, the model recovers the emphasis on Allen ($P = 1.0000$). On the factual version Allen drops to last place because the only edges remaining for him are routine professional and family ties — the evidence explicitly states his DNA exclusion, and the extractor reflects this with an innocent profile. Lawrence Kane rises to #1 because the facts surface three Kane-specific edges that no other suspect has: a workplace link to victim Donna Lass at the Sahara Tahoe casino in 1970, a witness identification from 1970 abduction survivor Kathleen Johns, and Donald Cheney’s testimony. Overall, the model has identified the current leading suspect as the primary villain.

3.2.6 Discussion

Where the model works. The strongest aggregate result is the precision/recall asymmetry: at 0.825 precision the R-GCN is a substantially more cautious detective than the feature-only LogReg, and the gain is paid for in 4 points of recall. The R-GCN solves cleanly (without falsely accusing anyone innocent) in 77.2% of test stories versus the LogReg’s 66.7%, and zero stories now go unsolved by both models. The disagreement structure points to the source of the gain: the R-GCN exclusively wins disagreements by clearing innocent individuals the LogReg flagged. Alibi, spatial separation, and social context only become useful evidence when the model can connect them across the graph, and the typed relations work for this. Table 5 confirms this is not a feature of graph topology in general: a Laplacian-eigenmaps baseline on a character-only weighted graph reaches 25.5% overall accuracy — essentially chance on a four-class task — and supplementing those embeddings onto the LogReg’s features adds nothing. The 32-point gap between the LogReg and the R-GCN is bought specifically by typed relational message passing.

The held-out Zodiac case study is the strongest qualitative result. Across two synopses of the same unsolved case — one a 2007 film, one a 2023–2025 factual record — the model produces different prime suspects, both consistent with the evidence in their respective sources, and re-ranks Arthur Leigh Allen from #1 ($P = 1.0000$) to #38 ($P = 0.0004$) when the synopsis foregrounds his DNA exclusion. We see this as evidence that the R-GCN is doing structural reading of the input graph.

Where the model fails. Recall is capped at roughly 0.71, and the cap is in the data, not the model. Feature analysis of the villains the R-GCN misses shows they have no villain-indicative features at all — no flagged motive, no concealment behavior, no hidden relationships. Of the originally-missed villains, 100% of caught villains had `has_motive`= 1 versus only 10% of missed villains; 78% of caught had `concealing_info`= 1 versus 10% of missed; 57% of caught had `hidden_relationship`= 1 versus 17.5% of missed. The extractor simply did not flag these traits, and no model architecture can recover signal that is not in the input. This was the dominant failure mode and the motivation for the three-round re-extraction loop, which closed most of the gap (villain F_1 from 0.703 to 0.762). As any good detective would do, the next logic step would be to collect more evidence in order to solve the crime.

Multi-villain ensembles are a second failure category. Recall degrades with the number of true villains in a story: 100% on single-villain cases, ~88–91% on two-villain cases, and roughly 39% on six-villain cases such as *Murder on the Orient Express*, where collective-guilt plots break the genre structure. Both the R-GCN and the LogReg fail here equivalently, which suggests it is a feature-distribution problem rather than a graph-structure problem: when “everyone” did it, the features that distinguish villain from non-villain in conventional whodunits no longer apply. We do not view this as a gap the architecture should be expected to close.

A third failure relates to mismatches the model could not have been expected to handle: TVE.089

(*Vera: Hidden Depths*), whose extracted villain entity was the literal phrase “People with specific forms of desperation and violence” rather than a named character, and TVE_093 (*Spiral: Series 4*), an organized-crime serial whose villain is a network rather than an individual. We excluded both and report the exclusions here.

4 Conclusion

This paper set out to survey the literature of Graph Autoencoders and then utilize a language model to produce graphs for murder mysteries for a Graph Convolutional Network to analyze, with the hope of it being able to predict villains. In the literature survey, we found a variety of methods, from classical spectral methods, to modern generative methods, that have been applied to graph learning. Our experiments with language model graph generation was fruitful, with the model able to produce accurate graphs for hundreds of murder mysteries. Finally, the Relational Graph Convolutional Network outperformed baselines in link prediction and node classification. This work is significant for three reasons: First, it utilized a language model to produce the training data. This project would have required many human hours just to produce the graphs, ignoring that this would require knowing the plots of hundreds of stories. This has the potential to impact fields of ML where the collection and pre-processing of text is a limiting factor. Second, it demonstrated that R-GCNs are capable of learning the complex relations present in relatively messy “real world” data. Our dataset was not a refined benchmark, and included various complications that the model was able to successfully handle. Finally, it was able to generalize from fake murder mystery stories into a real unsolved case. While not a significant result on its own, it is a sign that similar models could have applications to real world problems. Nonetheless, there were various limitations to our methods.

Finally, there are various extensions that we think warrant future attention. First, incorporating latent regularization into the R-GCN. This was not explored on the murder mystery dataset due to time limitations. Second, using a larger language model to gather more information about each mystery (if multi-modal language models advance, one could generate a graph from the first 80a movie). Third, there are plenty of hyperparameters that we did not explore due to compute and time limitations. Overall, the application of Graph Autoencoders to murder mysteries was a success, and we expect the field to be bright with the recent ability to leverage language models for graph generation.

References

- Uri Alon and Eran Yahav. On the bottleneck of graph neural networks and its practical implications. In *International Conference on Learning Representations*, 2021.
- Mikhail Belkin and Partha Niyogi. Laplacian eigenmaps and spectral techniques for embedding and clustering. In *Advances in Neural Information Processing Systems*, volume 14, 2001.
- Mikhail Belkin, Partha Niyogi, and Vikas Sindhwani. Manifold regularization: A geometric framework for learning from labeled and unlabeled examples. *Journal of Machine Learning Research*, 7: 2399–2434, 2006.
- Joan Bruna, Wojciech Zaremba, Arthur Szlam, and Yann LeCun. Spectral networks and locally connected networks on graphs. In *International Conference on Learning Representations*, 2014.

- Michaël Defferrard, Xavier Bresson, and Pierre Vandergheynst. Convolutional neural networks on graphs with fast localized spectral filtering. In *Advances in Neural Information Processing Systems*, volume 29, 2016.
- Vijay Prakash Dwivedi and Xavier Bresson. A generalization of transformer networks to graphs. *arXiv preprint arXiv:2012.09699*, 2020.
- Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1263–1272. PMLR, 2017.
- Aditya Grover, Aaron Zweig, and Stefano Ermon. Graphite: Iterative generative modeling of graphs. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2434–2444. PMLR, 2019.
- Ziniu Hu, Yuxiao Dong, Kuansan Wang, and Yizhou Sun. Heterogeneous graph transformer. In *Proceedings of The Web Conference 2020*, pages 2704–2710. Association for Computing Machinery, 2020. doi: 10.1145/3366423.3380027.
- Diederik P. Kingma and Max Welling. Auto-encoding variational bayes. In *International Conference on Learning Representations*, 2014.
- Thomas N. Kipf and Max Welling. Variational graph auto-encoders. *arXiv preprint arXiv:1611.07308*, 2016.
- Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations*, 2017.
- Qimai Li, Zhichao Han, and Xiao-Ming Wu. Deeper insights into graph convolutional networks for semi-supervised learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018.
- Moritz Plenz and Anette Frank. Graph language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4477–4494, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.245.
- Michael Schlichtkrull, Thomas N. Kipf, Peter Bloem, Rianne van den Berg, Ivan Titov, and Max Welling. Modeling relational data with graph convolutional networks. In *The Semantic Web*, volume 10843 of *Lecture Notes in Computer Science*, pages 593–607. Springer, 2018. doi: 10.1007/978-3-319-93417-4_38.
- David I. Shuman, Sunil K. Narang, Pascal Frossard, Antonio Ortega, and Pierre Vandergheynst. The emerging field of signal processing on graphs: Extending high-dimensional data analysis to networks and other irregular domains. *IEEE Signal Processing Magazine*, 30(3):83–98, 2013. doi: 10.1109/MSP.2012.2235192.
- Martin Simonovsky and Nikos Komodakis. GraphVAE: Towards generation of small graphs using variational autoencoders. In *Artificial Neural Networks and Machine Learning – ICANN 2018*, volume 11139 of *Lecture Notes in Computer Science*, pages 412–422. Springer, 2018. doi: 10.1007/978-3-030-01418-6_41.

- Petar Veličković. Message passing all the way up. In *ICLR 2022 Workshop on Geometrical and Topological Representation Learning*, 2022.
- Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua Bengio. Graph attention networks. In *International Conference on Learning Representations*, 2018.
- Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? In *International Conference on Learning Representations*, 2019.
- Chengxuan Ying, Tianle Cai, Shengjie Luo, Shuxin Zheng, Guolin Ke, Di He, Yanming Shen, and Tie-Yan Liu. Do transformers really perform bad for graph representation? In *Advances in Neural Information Processing Systems*, volume 34, pages 28877–28888, 2021.